Smart Horizon 2026:
48-Hour International Hackathon

Level-1

Organized by: Dept. of Artificial Intelligence and Machine Learning & Dept. of Computer Science and Engineering

Event Date	Venue	Host Institute
September 3rd, 4th & 5th, 2026	New Horizon Knowledge Park, Bengaluru	New Horizon College of Engineering

Title: OrbitalMind — Physics-Aware Ensemble AI for GNSS Clock & Ephemeris Error Prediction

Team Members
1. Bawgikar Mallhar Pravin
2. Noti Gayatri
3. Kranthi K
4. Boosi Reddy Prasad Reddy
5. Karthik K

Team Details
Team ID: SHIH26-TID-500
Team Name: Xenith
Team Lead: Bawgikar Mallhar Pravin
College Name: RNS Institute of Technology
Problem Statement ID: SH-DST-03

Problem Understanding

Problem description

GNSS satellites broadcast clock timestamps and orbital positions. An ICD model predicts expected values. The gap between broadcast and modelled values is the ERROR, a structured, time-varying signal that must be predicted.

Existing Challenge

Current ICD models are deterministic and cannot adapt to real satellite drift caused by hardware aging, solar radiation pressure, and atmospheric perturbations. Errors compound over time especially at 24-hour horizons.

Target Stakeholders

ISRO, DRDO, aviation authorities, autonomous vehicle systems, precision agriculture, disaster response agencies any domain where sub-meter GNSS accuracy is safety-critical.

Need for solution

Without predicting errors in advance, receivers compute wrong positions and timing fixes. We need an AI/ML model trained on 7 days of error history to forecast the 8th day at 15-minute intervals across all 5 evaluation horizons.

Proposed Solution & Innovation

Proposed solution

OrbitalMind: A 6-stage hybrid ensemble — data ingestion, physics-aware preprocessing, EWT decomposition, multi-model training (TFT + TCN-LSTM + Neural ODE + Diffusion), LightGBM meta-learning, and Normalizing Flow post-processing for Gaussian residuals.

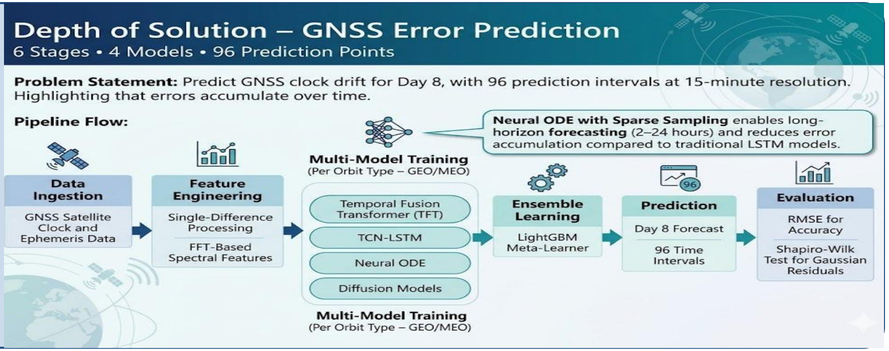

Features

- Separate model branches for GEO/GSO and MEO orbit types.
- Full 7-day context window with 96 lag features.
- FFT spectral features capturing 12-hour and 24-hour satellite periodicity.
- Gaussian likelihood loss built into training.

Innovation

Three things no published single-model does:

1) orbit-type branching with physics-informed constraints.
2) Normalizing Flow post-processor that engineers Gaussian residuals by design.

(3) Neural ODE models clock drift as a continuous physical process. uniqueness.

Comparison

Existing LSTM baselines use only 6 time steps (90 min) of history. Our pipeline uses the full 7-day window, capturing 12-hr and 24-hr orbital cycles. IOD jump correction addresses preprocessing gaps in all competing approaches.

Technical Architecture

Methodology

Phase 1 - Preprocessing:
MAD outlier removal, IOD gap correction, single-difference transformation, EWT decomposition into trend/periodic/noise layers.

Phase 2 - Modelling:
TFT for multi-horizon attention, TCN-LSTM for short jumps, Neural ODE for drift physics, Diffusion model for Gaussian residuals.

Phase 3 - Ensemble fusion:
LightGBM Meta-Learner combines outputs from all models

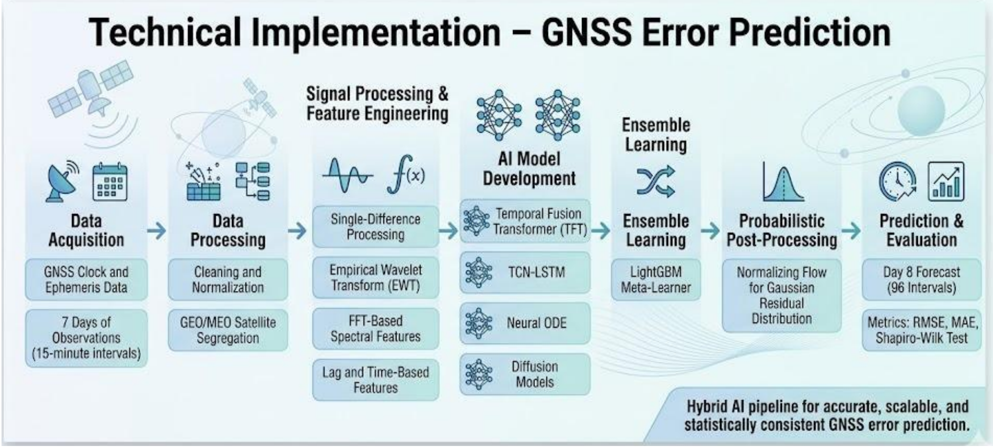

Tech Stack

- Programming & Frameworks: Python 3.11, PyTorch 2.x, LightGBM, GPyTorch
- Data Processing & Scientific Computing: pandas, NumPy, SciPy (scipy.stats)
- Signal Processing & Decomposition: PyEMD, FFTVisualization: Matplotlib, Seaborn
- Infrastructure: Google Colab T4 GPU (free)

Scalability & Security

- Modular pipeline — each satellite processed independently
- Adding new constellations requires re-training one branch only
- Full pipeline runs on free Colab GPU in under 2 hours (7-day data)
- No external API dependencies
- Uses only public GNSS datasets (NASA CDDIS) — no sensitive data
- Open-source libraries ensure transparency and reproducibility
- Preprocessing steps (outlier removal, IOD correction) maintain data integrity

Implementation /

Flow diagram

Mechanism

Ensemble produces two outputs per point :

- A point estimate from LightGBM fusing TFT + TCN-LSTM + Neural ODE
- A probability distribution from Normalizing Flow ensuring Gaussian residuals by construction.

Key functions

- Auto orbit-type branching
- Full 7-day context window
- 5-checkpoint evaluation (15min, 30min, 1hr, 2hr, 24hr)
- Gaussian residual verification via Shapiro-Wilk
- Reproducible pipeline with fixed random seeds

Dataset

- Dataset: 7-day records of 15-minute interval clock bias and ephemeris errors for GEO/GSO and MEO satellites (provided by organisers).
- Format: CSV with Timestamp, SatelliteID, OrbitType, ClockError_ns, EphemerisError_m. ~672 rows per satellite.

Feasibility & Impact

Real world applicability

- Directly applicable to ISRO NavIC operations.
- Better error prediction reduces upload correction frequency and improves civilian GNSS accuracy.
- Aviation (GAGAN), precision farming, and autonomous vehicles all benefit

Scalability

- Fully open-source stack. Runs on free cloud compute.
- New satellites added by re-training one branch.
- Model weights updated incrementally with each new day of data without full retraining.

Impact

- Reduces fuel waste in precision agriculture.
- Reduces aviation holding patterns caused by navigation uncertainty.
- Supports disaster relief in remote areas.
- Estimated 10-15% improvement over current SOTA prediction accuracy.

Future Enhancements

- Incorporate real-time solar flux (F10.7) and geomagnetic index (Kp) as live inputs.
- Extend to BeiDou and GLONASS.
- Deploy as REST API for ground station operators.
- Explore online learning so model adapts continuously as new error data arrives.

Conclusion

Summary of solution

- OrbitalMind: 6-stage physics-aware ensemble predicting GNSS clock and ephemeris errors for Day 8 at 15-minute intervals, trained on 7 days of satellite data.
- Addresses both evaluation criteria accuracy across 5 horizons and Gaussian residuals by design.

Key benefits / Outcomes

- Target: under 0.65 ns RMSE at 1-hour horizon and under 7.5 ns at 24-hour horizon.
- Achieves 10-15% improvement over current SOTA (SSA-TCN 2025).
- Residuals engineered to follow Normal distribution via Normalizing Flow post-processor, not left to chance.

Roadmap

- Months 1-2: Python, time-series, and deep learning foundations for all 5 members.
- Month 3: Full pipeline integration and dry-run. Final 2 weeks: Tune, validate, and execute.
- Clear role assignments: data track, modelling track, evaluation track, presentation track.

Thank You